KOOKMIN UNIVERSITY

BentoML

소프트웨어융합대학원
진혜진

실습(Hugging Face Transformer 모델)

- 가상환경 활성화
  - python -m venv venv
  - venv\Scripts\activate

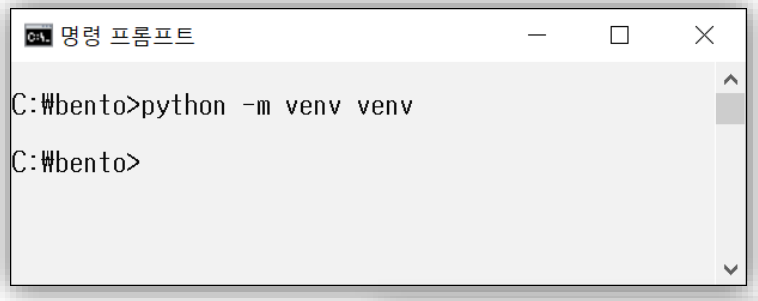

- 종속성 설치
  - pip install bentoml torch transformers

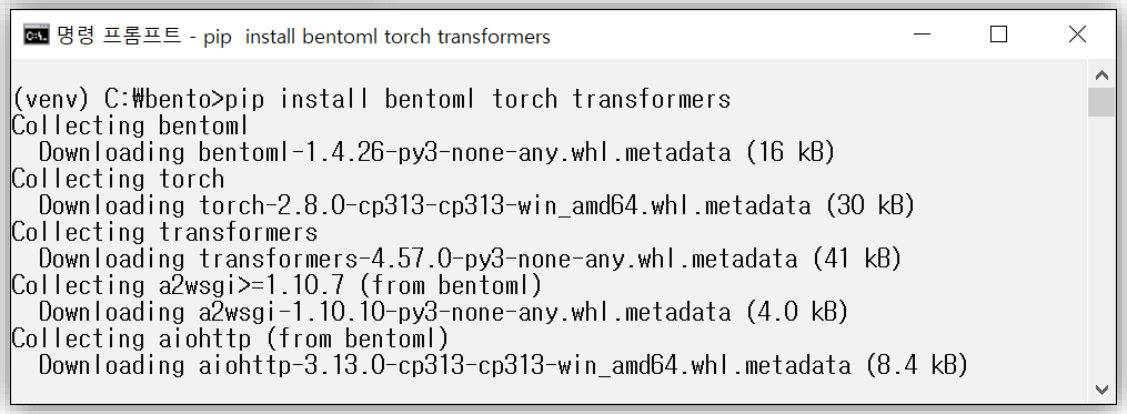

- 설치 확인
  - pip show bentoml
  - pip show torch
  - pip show transformers

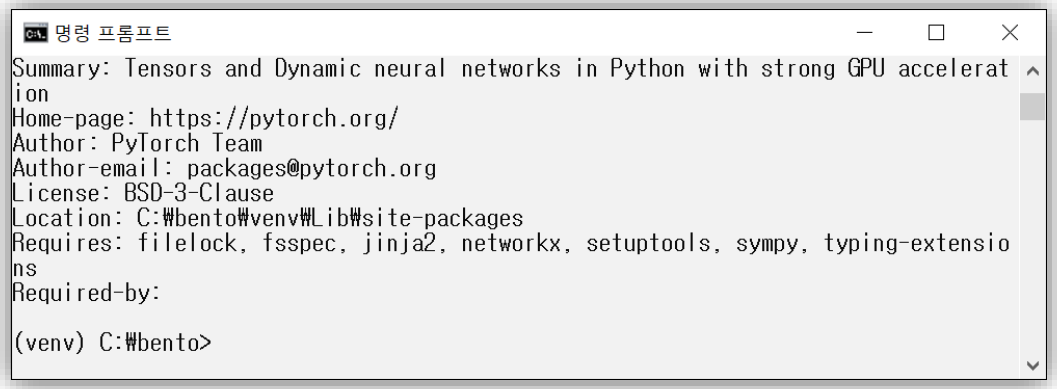

서비스(API) 만들기

- service.py

```python
from __future__ import annotations
import bentoml

with bentoml.importing():
    from transformers import pipeline


EXAMPLE_INPUT = "This is a demonstration of BentoML with a summarization model."


@bentoml.service
class Summarization:
    def __init__(self) -> None:
        self.pipeline = pipeline("summarization")

    @bentoml.api
    def summarize(self, text: str = EXAMPLE_INPUT) -> str:
        result = self.pipeline(text)
        return f"Here's your summary: {result[0]['summary_text']}"
```

- 서버 실행
  - bentoml serve
- http://localhost:3000

명령 프롬프트 - bentoml serve

```
Summary: Tensors and Dynamic neural networks in Python with strong GPU acceleration
Home-page: https://pytorch.org/
Author: PyTorch Team
Author-email: packages@pytorch.org
License: BSD-3-Clause
Location: C:₩bento₩venv₩Lib₩site-packages
Requires: filelock, fsspec, jinja2, networkx, setuptools, sympy, typing-extensions
Required-by:

(venv) C:₩bento>bentoml serve
2025-10-11T10:20:14+0900 [INFO] [cli] Loading service from default location 'service.py'
2025-10-11T10:20:19+0900 [INFO] [cli] Loading service from default location 'service.py'
2025-10-11T10:20:20+0900 [INFO] [cli] Starting production HTTP BentoServer from "." listening on http://localhost:3000 (Press CTRL+C
 to quit)
2025-10-11T10:20:21+0900 [INFO] [:1] Loading service from default location 'service.py'
No model was supplied, defaulted to sshleifer/distilbart-cnn-12-6 and revision a4f8f3e (https://huggingface.co/sshleifer/distilbart-
cnn-12-6).
Using a pipeline without specifying a model name and revision in production is not recommended.
```

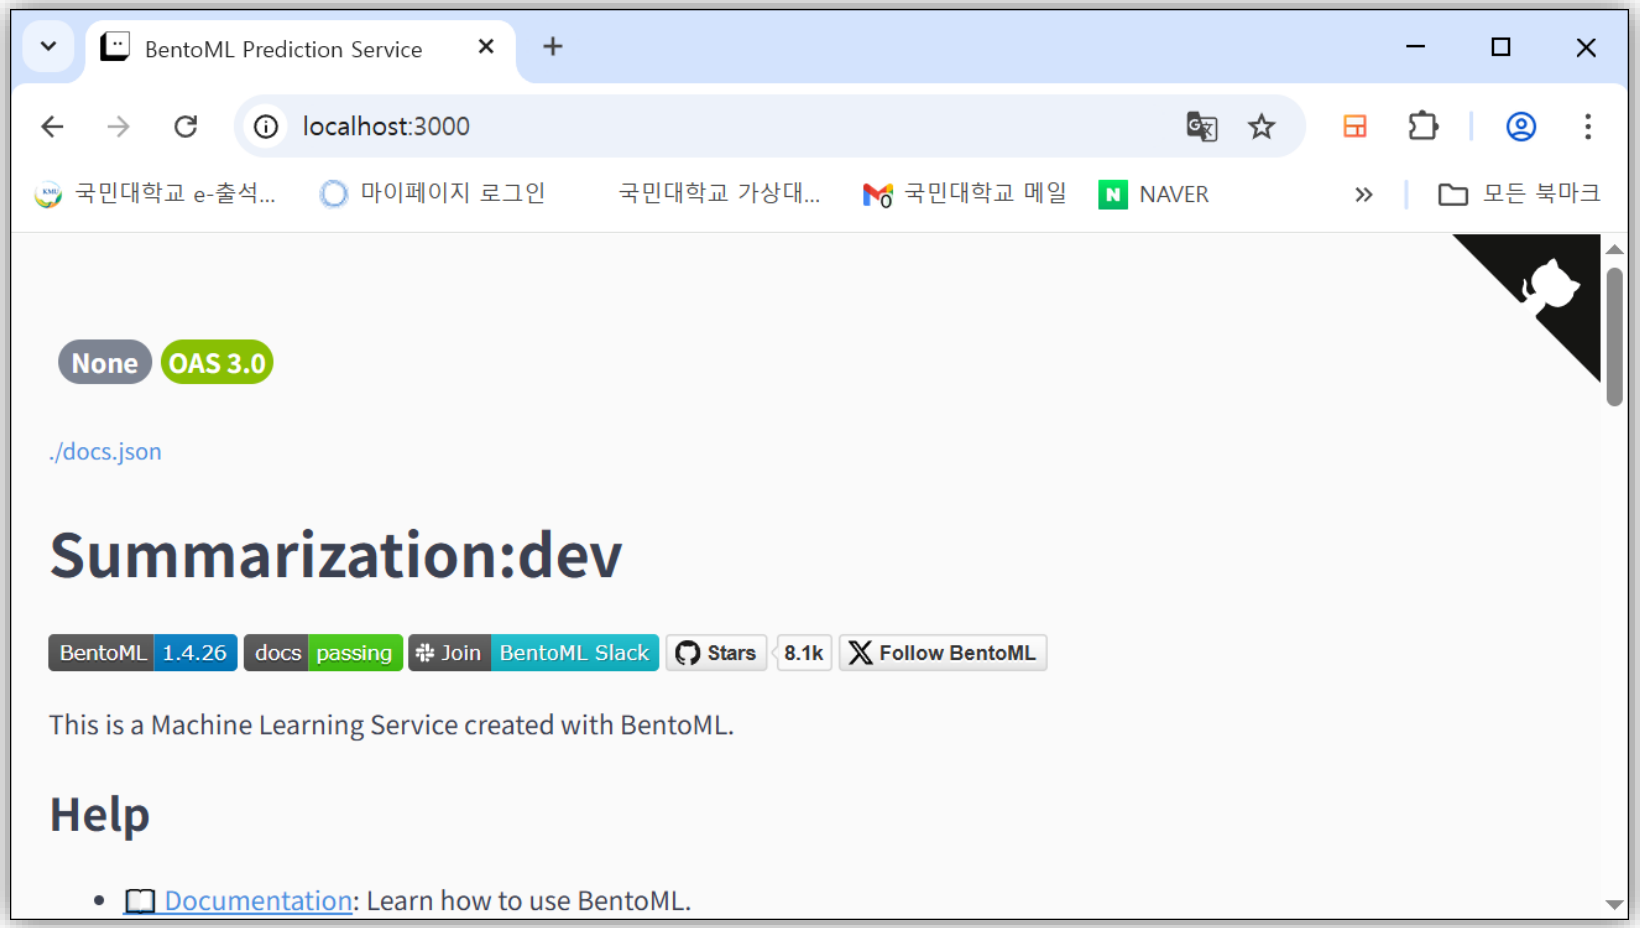
BentoML Prediction Service
localhost:3000
국민대학교 e-출석...
마이페이지 로그인
국민대학교 가상대...
국민대학교 메일
NAVER
모든 북마크
None
OAS 3.0
./docs.json
Summarization:dev
BentoML 1.4.26
docs passing
Join BentoML Slack
Stars 8.1k
Follow BentoML
This is a Machine Learning Service created with BentoML.
Help
Documentation: Learn how to use BentoML.

- 요청 보내기, 응답 확인(Swagger UI)
  - http://localhost:3000
  - Service APIs - summarize - Try it out

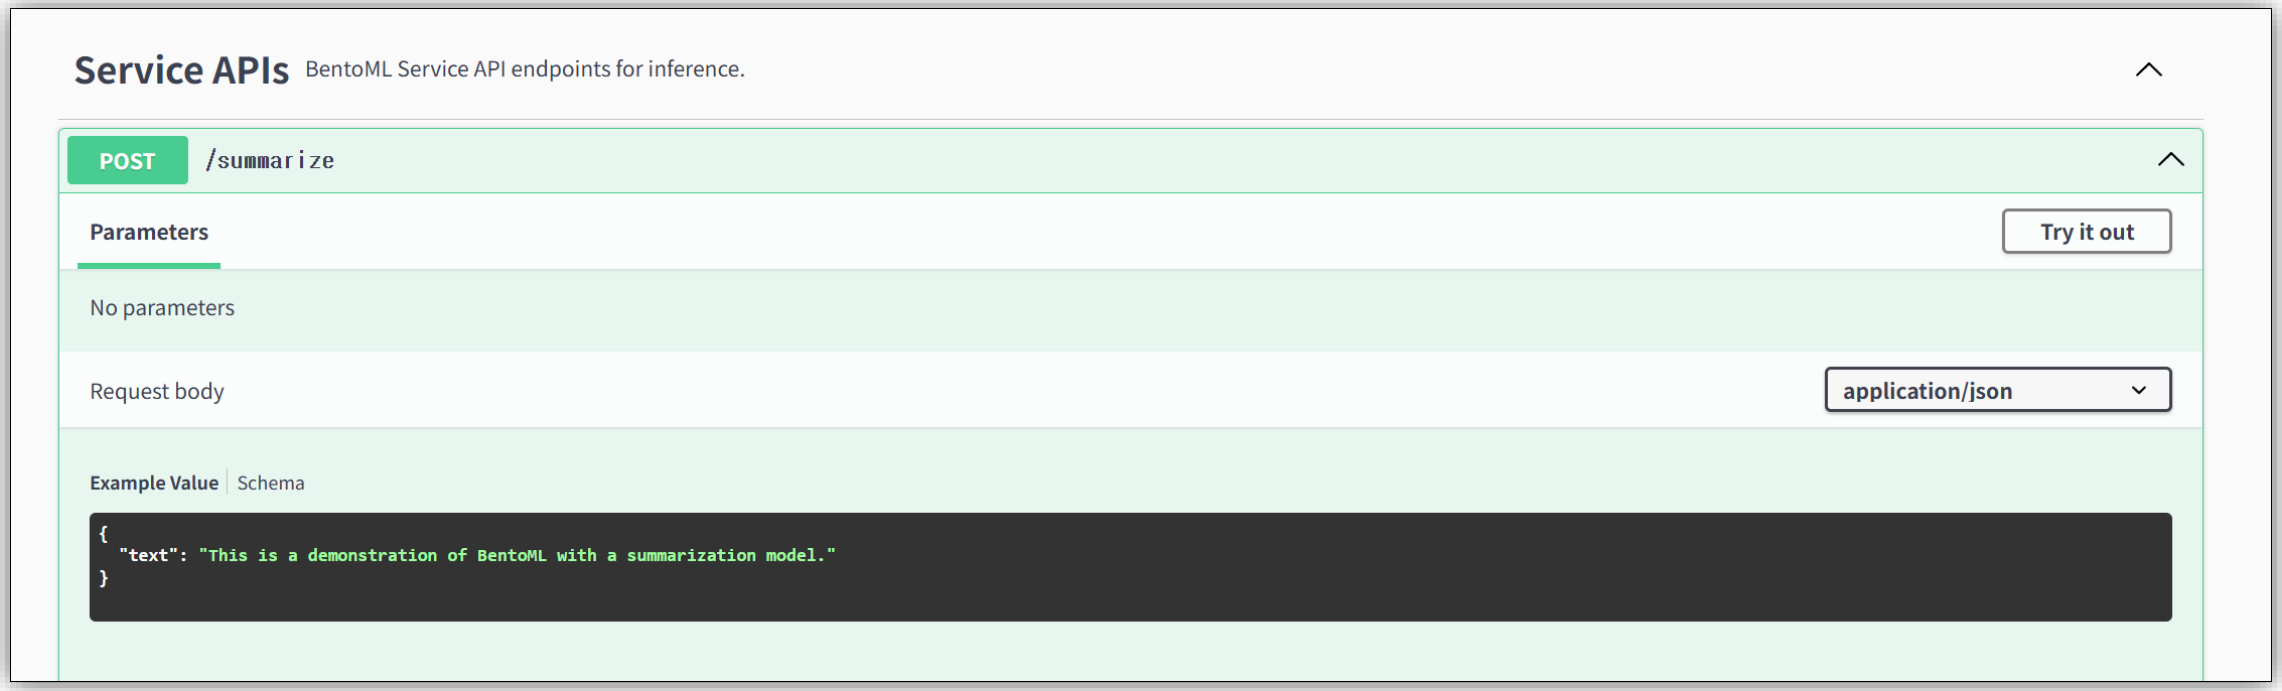

Service APIs BentoML Service API endpoints for inference.

POST /summarize

Parameters Cancel

No parameters

Request body application/json

```
{
  "text": "This is a demonstration of BentoML with a summarization model."
}
```

Execute

Responses

Curl

```
curl -X 'POST' \
  'http://localhost:3000/summarize' \
  -H 'accept: text/plain' \
  -H 'Content-Type: application/json' \
  -d '{
  "text": "This is a demonstration of BentoML with a summarization model."
}'
```

Request URL

```
http://localhost:3000/summarize
```

Server response

Code	Details
200	Response body

```
Here's your summary:  This is a demonstration of BentoML with a summarization model . This is the first time Bento ML has been shown to be able to work with a large number of users . The model i
s based on a model of a search for a user's search engine . The search engine is now being used to analyse the results .
```

Download

Response headers

```
content-length: 314
content-type: text/plain; charset=utf-8
date: Sat,11 Oct 2025 01:32:22 GMT
server: BentoML Service/Summarization
x-bentoml-request-id: 5312a53011799897
```

실습(scikit-learn Iris 분류 모델)

- train_model.py

```python
from sklearn import datasets
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
import bentoml

iris = datasets.load_iris()
X_train, X_test, y_train, y_test = train_test_split(
    iris.data, iris.target, test_size=0.2, random_state=42
)


model = RandomForestClassifier()
model.fit(X_train, y_train)


bentoml.sklearn.save_model(
    "iris_rf_model",
    model,
    signatures={"predict": {"batchable": True}},
)

print("✅ 혜진님 모델이 저장되었습니다!")
```

실행

- python train_model.py

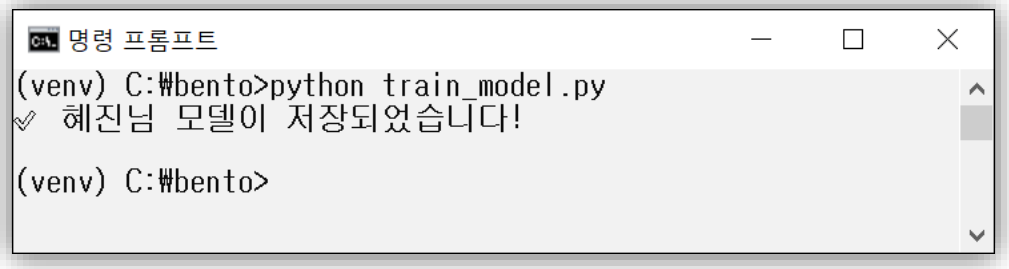

- pip install scikit-learn

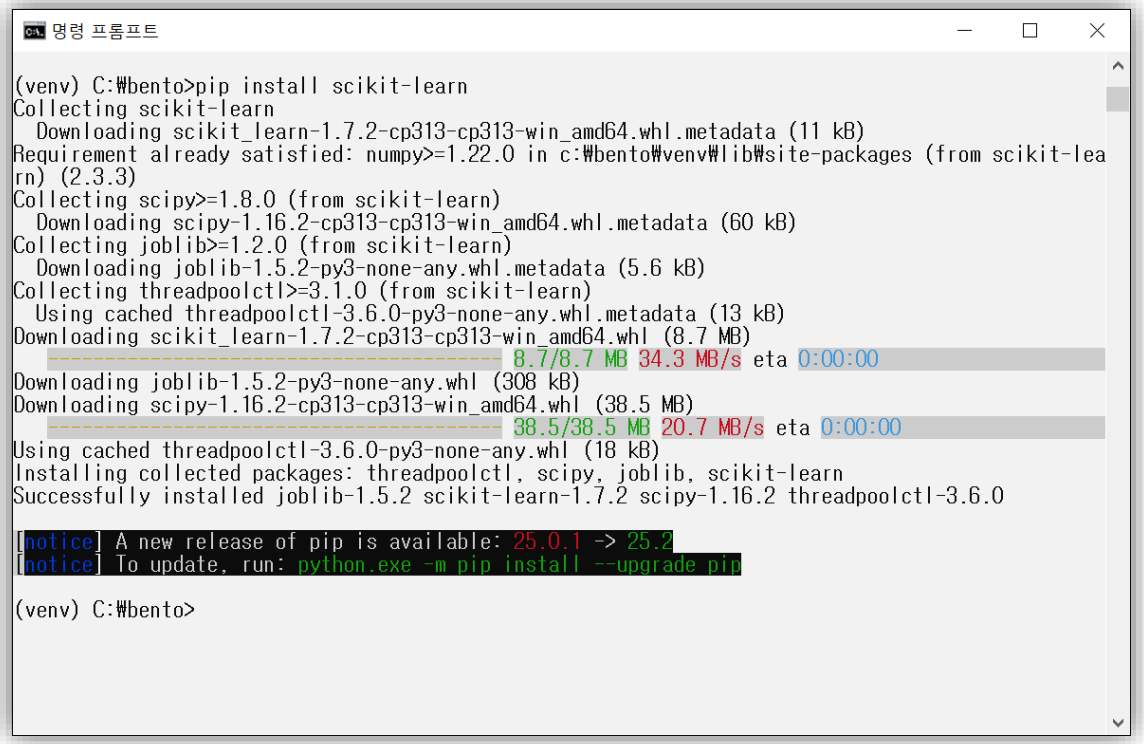

서비스(API) 만들기

```python
import bentoml
from pydantic import BaseModel
import numpy as np

# 요청/응답 스키마
class IrisInput(BaseModel):
    sepal_length: float
    sepal_width: float
    petal_length: float
    petal_width: float

class IrisOutput(BaseModel):
    prediction: int

# Service 정의 (클래스 기반)
@bentoml.service(name="iris_classifier")
class IrisService:
    def __init__(self):

        self.model = bentoml.models.get("iris_rf_model:latest").load_model()

    @bentoml.api
    def predict(self, data: IrisInput) -> IrisOutput:
        features = np.array([[
            data.sepal_length, data.sepal_width,
            data.petal_length, data.petal_width
        ]])
        pred = self.model.predict(features)
        return IrisOutput(prediction=int(pred[0]))
```

서비스 실행

명령 프롬프트 - bentoml serve service.py:IrisService

```
(venv) C:₩bento>bentoml serve service.py:IrisService
2025-10-11T11:30:38+0900 [INFO] [cli] Starting production HTTP BentoServer from "service.py:IrisService" listen
ing on http://localhost:3000 (Press CTRL+C to quit)
2025-10-11T11:30:40+0900 [INFO] [entry_service:iris_classifier:1] Service iris_classifier initialized
```

BentoML Prediction Service

localhost:3000

국민대학교 e-출석... 마이페이지 로그인 국민대학교 가상대... 국민대학교 메일 NAVER 모든 북마크

None OAS 3.0

./docs.json

iris_classifier:dev

BentoML 1.4.26 docs passing Join BentoML Slack Stars 8.1k Follow BentoML

This is a Machine Learning Service created with BentoML.

Help

- Documentation: Learn how to use BentoML.

테스트

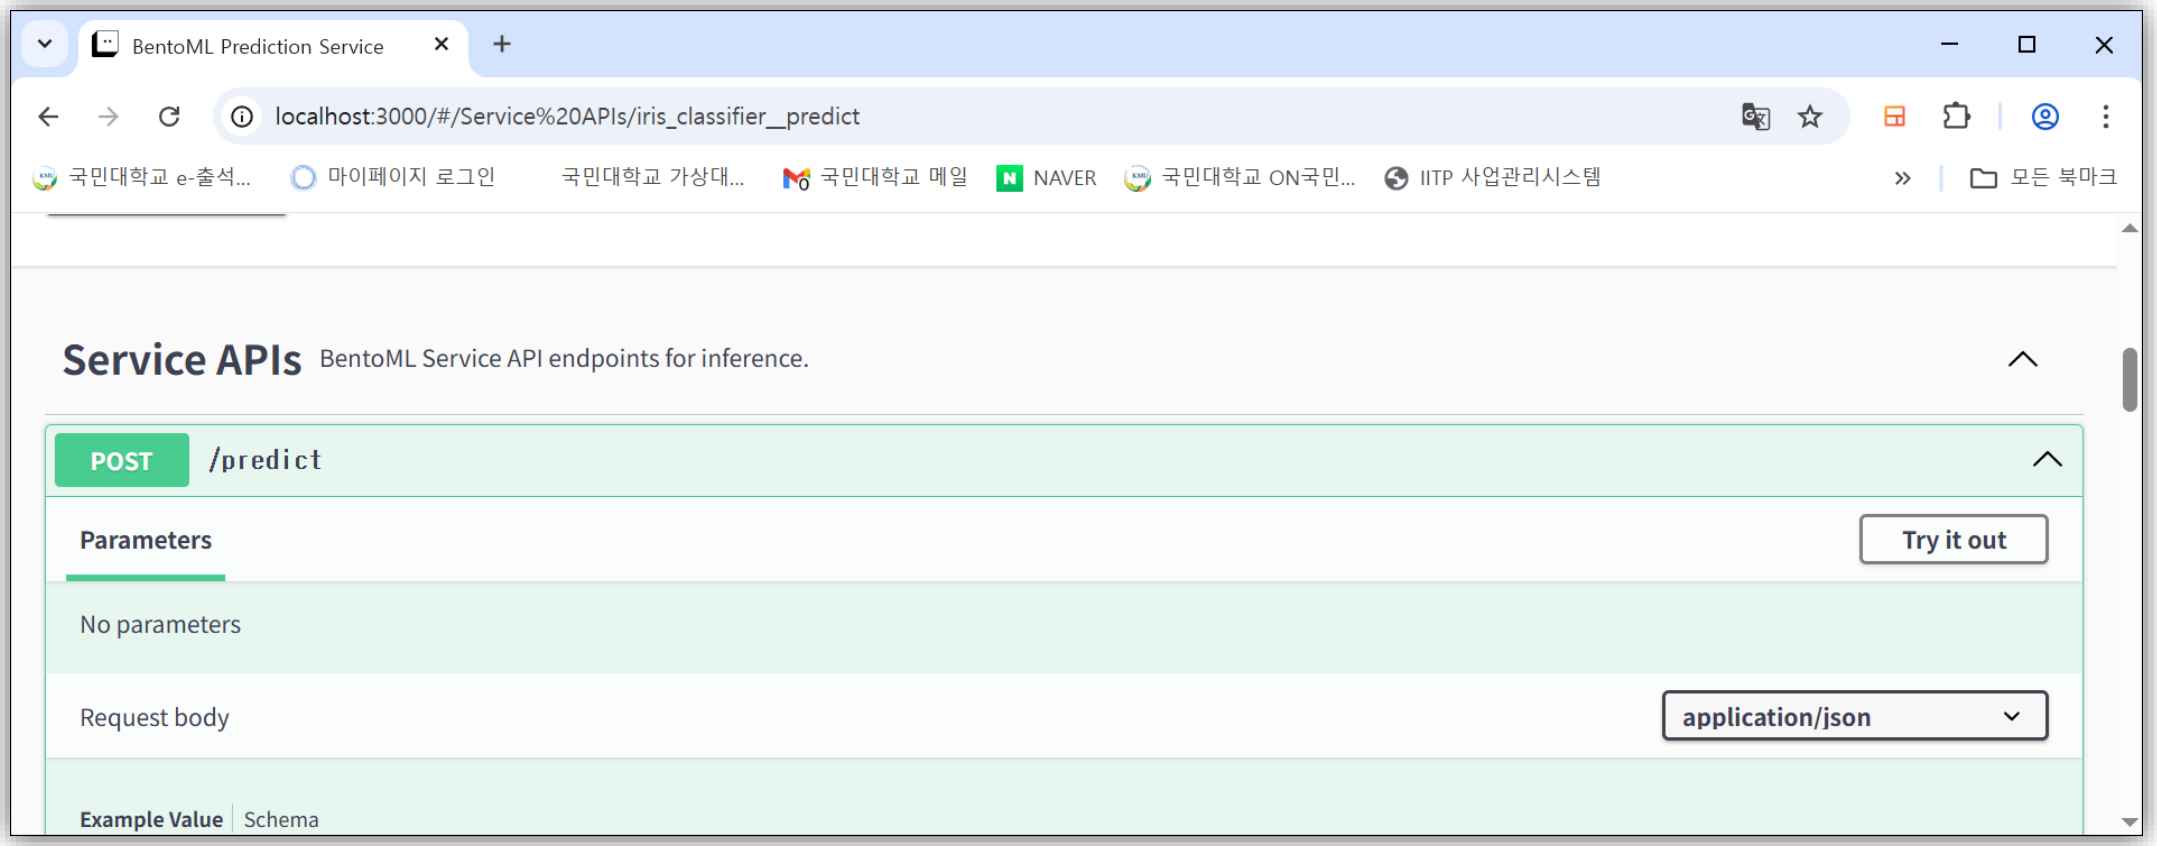
BentoML Prediction Service
localhost:3000/#/Service%20APIs/iris_classifier__predict
국민대학교 e-출석...
마이페이지 로그인
국민대학교 가상대...
국민대학교 메일
NAVER
국민대학교 ON국민...
IITP 사업관리시스템
모든 북마크
Service APIs
BentoML Service API endpoints for inference.
POST
/predict
Parameters
Try it out
No parameters
Request body
application/json
Example Value
Schema

Request body

```
{
  "data": {
  "sepal_length": 5.1,
  "sepal_width": 3.5,
  "petal_length": 1.4,
  "petal_width": 0.2
  }
}
```

Execute

Responses

Curl

```
curl -X 'POST' \
  'http://localhost:3000/predict' \
  -H 'accept: application/json' \
  -H 'Content-Type: application/json' \
  -d '{
  "data": {
  "sepal_length": 5.1,
  "sepal_width": 3.5,
  "petal_length": 1.4,
  "petal_width": 0.2
  }
}'
```

Request URL

```
http://localhost:3000/predict
```

Server response

Code	Details
200	Response body

```
{
  "prediction": 0
}
```

Server response

Code	Details
200	Response body

```
{
  "prediction": 1
}
```

Response headers

```
content-length: 16
content-type: application/json
date: Sat,11 Oct 2025 02:38:22 GMT
server: BentoML Service/iris_classifier
x-bentoml-request-id: 0c7aa0e2588c62e7
```

Responses